

Name : Arun M Register No : 23122110 Class : 3MScDS B

#### A. Data Exploration and Preprocessing

```
from sklearn.datasets import fetch_california_housing
import pandas as pd

# Load the California housing dataset
housing = fetch_california_housing()

# Create a DataFrame from the dataset using the data and feature names
df = pd.DataFrame(data=housing.data, columns=housing.feature_names)

# Print the DataFrame to view the data
print(df)

# Print the keys of the housing dataset to see available information
print(housing.keys())

dict_keys(['data', 'target', 'frame', 'target_names', 'feature_names', 'DESCR'])

#Finding the First five Rows
df.head()
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude \						
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556
37.88						
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842
37.86						
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260
37.85						
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945
37.85						
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467
37.85						

	Longitude	House Price
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

```
#Finding the descriptive Statistics of the All Variables
df.describe()
```

	MedInc	HouseAge	AveRooms	AveBedrms
Population \				
count	20640.000000	20640.000000	20640.000000	20640.000000

```

20640.000000
mean      3.870671      28.639486      5.429000      1.096675
1425.476744
std       1.899822      12.585558      2.474173      0.473911
1132.462122
min       0.499900      1.000000      0.846154      0.333333
3.000000
25%       2.563400      18.000000      4.440716      1.006079
787.000000
50%       3.534800      29.000000      5.229129      1.048780
1166.000000
75%       4.743250      37.000000      6.052381      1.099526
1725.000000
max       15.000100      52.000000      141.909091      34.066667
35682.000000

```

```

          AveOccup      Latitude      Longitude
count  20640.000000  20640.000000  20640.000000
mean      3.070655      35.631861    -119.569704
std      10.386050      2.135952      2.003532
min       0.692308      32.540000    -124.350000
25%       2.429741      33.930000    -121.800000
50%       2.818116      34.260000    -118.490000
75%       3.282261      37.710000    -118.010000
max      1243.333333      41.950000    -114.310000

```

### *#DataTypes of the Dataset*

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   MedInc      20640 non-null  float64
 1   HouseAge    20640 non-null  float64
 2   AveRooms    20640 non-null  float64
 3   AveBedrms   20640 non-null  float64
 4   Population  20640 non-null  float64
 5   AveOccup    20640 non-null  float64
 6   Latitude    20640 non-null  float64
 7   Longitude   20640 non-null  float64
dtypes: float64(8)
memory usage: 1.3 MB

```

### *#Descriptive Statistics of All Features including all Categorical Variables*

```
df.describe(include="all")
```

	MedInc	HouseAge	AveRooms	AveBedrms
Population \				
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675
std	1.899822	12.585558	2.474173	0.473911
min	0.499900	1.000000	0.846154	0.333333
25%	2.563400	18.000000	4.440716	1.006079
50%	3.534800	29.000000	5.229129	1.048780
75%	4.743250	37.000000	6.052381	1.099526
max	15.000100	52.000000	141.909091	34.066667

	AveOccup	Latitude	Longitude
count	20640.000000	20640.000000	20640.000000
mean	3.070655	35.631861	-119.569704
std	10.386050	2.135952	2.003532
min	0.692308	32.540000	-124.350000
25%	2.429741	33.930000	-121.800000
50%	2.818116	34.260000	-118.490000
75%	3.282261	37.710000	-118.010000
max	1243.333333	41.950000	-114.310000

*#Adding the target variable to the Dataset*

```
df['House Price'] = housing.target
df
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude \						
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467
...	...	...	...	...	...	...
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807

39.49						
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635
39.43						
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209
39.43						
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981
39.37						

	Longitude	House Price
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422
...	...	...
20635	-121.09	0.781
20636	-121.21	0.771
20637	-121.22	0.923
20638	-121.32	0.847
20639	-121.24	0.894

[20640 rows x 9 columns]

*#Checking is there any Missing Values*

```
missing=df.dropna(inplace = True)
```

```
print(missing)
```

None

*#Checking for Duplicate Values*

```
dup=df.drop_duplicates(inplace = True)
```

```
print(dup)
```

None

*#Checking the Null Values in the Dataset*

```
null_counts = df.isnull().sum()
```

```
print("Null values in each column:")
```

```
print(null_counts)
```

Null values in each column:

MedInc 0

HouseAge 0

AveRooms 0

AveBedrms 0

Population 0

AveOccup 0

Latitude 0

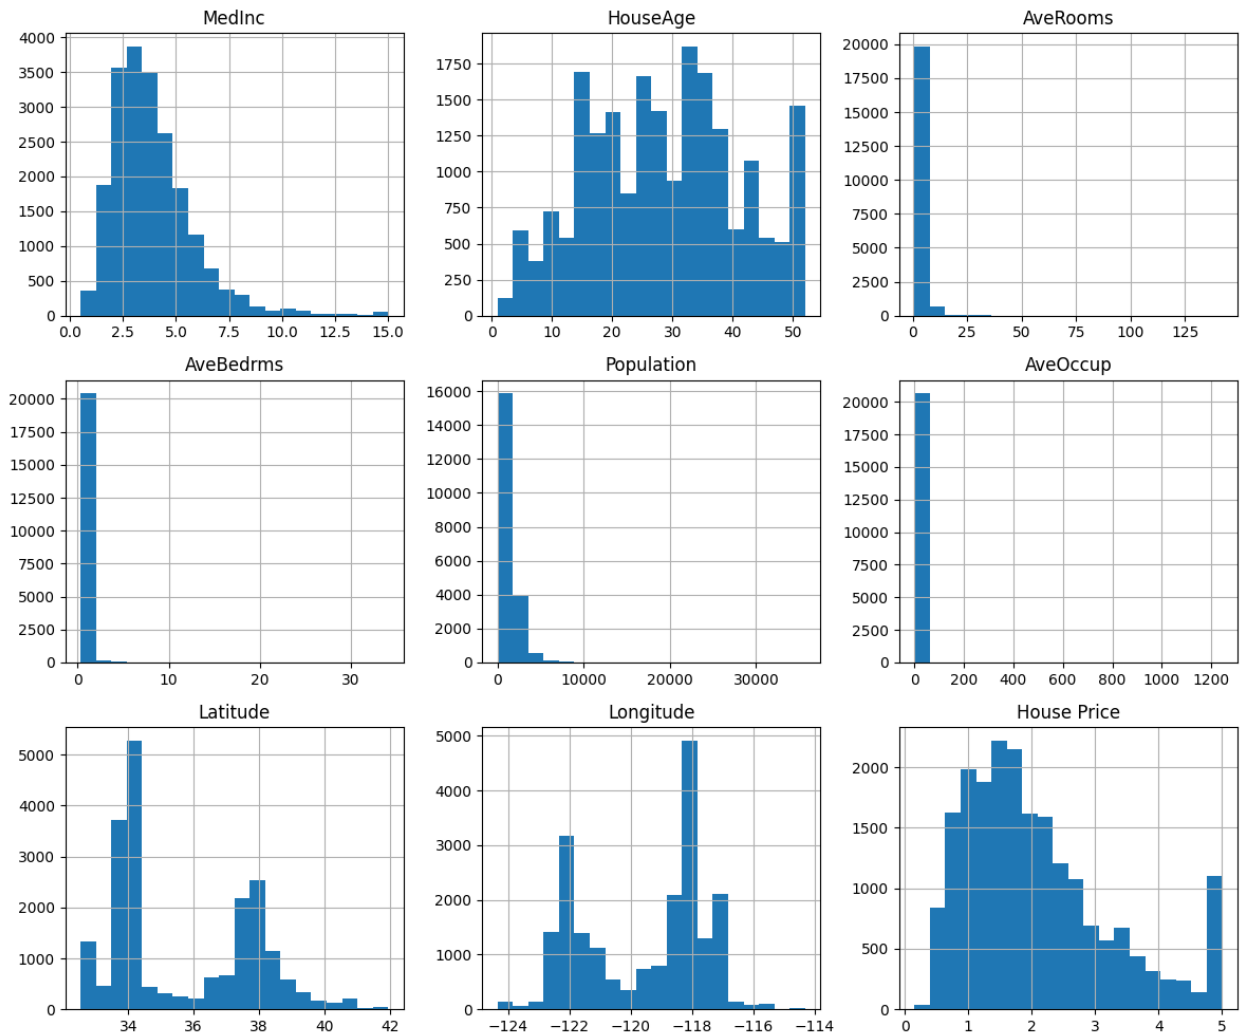
Longitude 0

House Price 0

dtype: int64

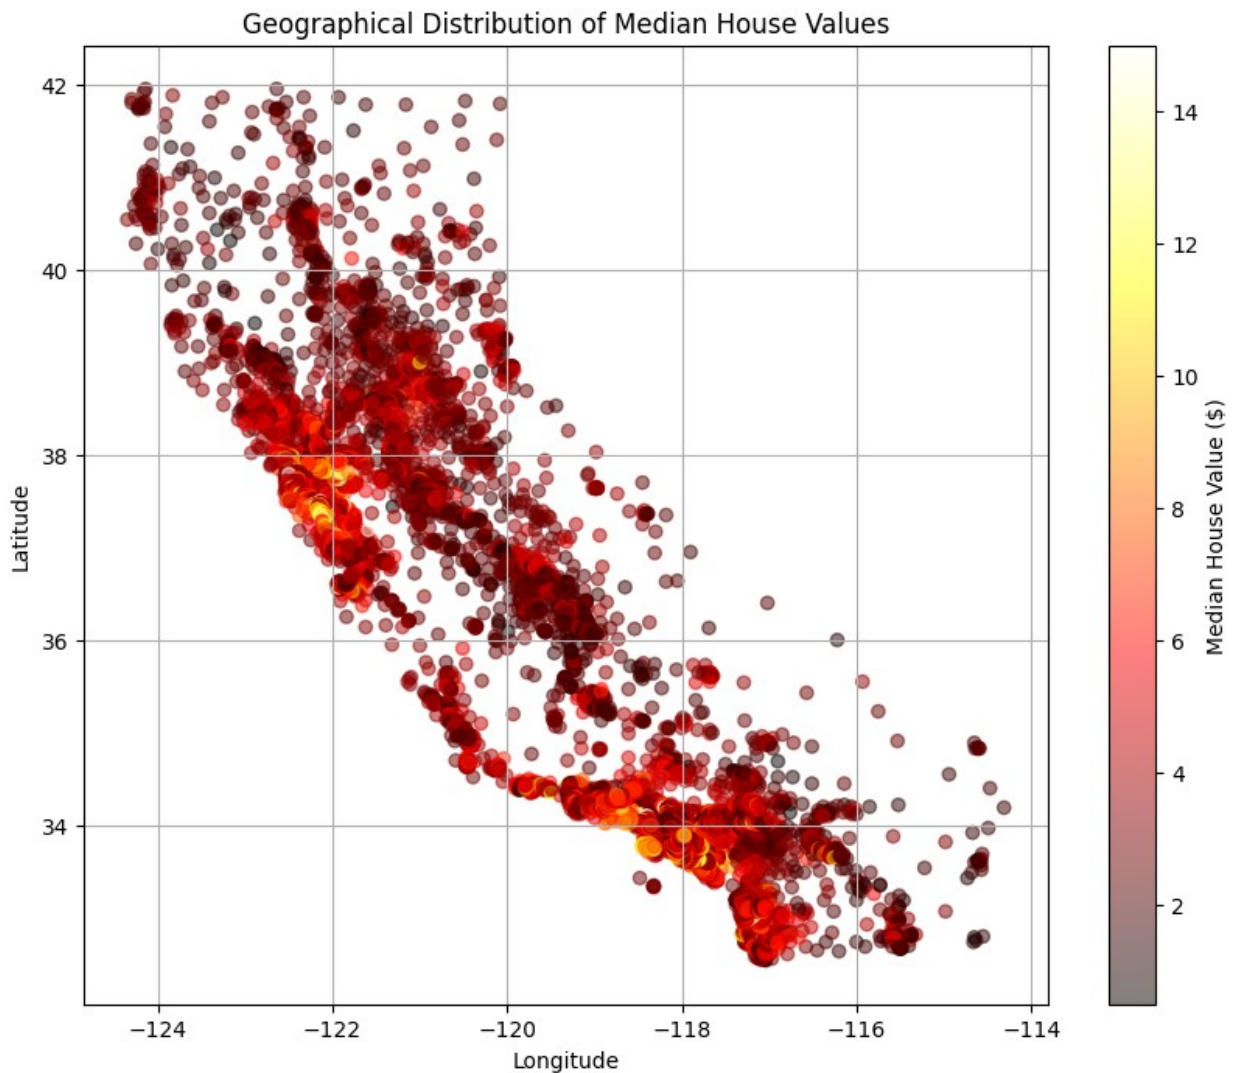
```
#Histogram of all Features
import matplotlib.pyplot as plt

df.hist(figsize=(12,10),bins=20)
plt.tight_layout()
plt.show()
```



```
#This is the Scatter Plot Geographical Distribution of Median House
Values
plt.figure(figsize=(10, 8))
plt.scatter(df['Longitude'], df['Latitude'], c=df['MedInc'],
cmap='hot', alpha=0.5)
plt.colorbar(label='Median House Value ($)')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.title('Geographical Distribution of Median House Values')
```

```
plt.grid(True)
plt.show()
```



## B. Model Building

*#Finding the Correlation to Extract the Best Input Variables*

```
corr=df.corr()
corr
```


	MedInc	HouseAge	AveRooms	AveBedrms	Population
Ave0ccup \					
MedInc	1.000000	-0.119034	0.326895	-0.062040	0.004834
0.018766					
HouseAge	-0.119034	1.000000	-0.153277	-0.077747	-0.296244
0.013191					
AveRooms	0.326895	-0.153277	1.000000	0.847621	-0.072213
0.004852					

AveBedrms	-0.062040	-0.077747	0.847621	1.000000	-0.066197	-
0.006181						
Population	0.004834	-0.296244	-0.072213	-0.066197	1.000000	
0.069863						
AveOccup	0.018766	0.013191	-0.004852	-0.006181	0.069863	
1.000000						
Latitude	-0.079809	0.011173	0.106389	0.069721	-0.108785	
0.002366						
Longitude	-0.015176	-0.108197	-0.027540	0.013344	0.099773	
0.002476						
House Price	0.688075	0.105623	0.151948	-0.046701	-0.024650	-
0.023737						

	Latitude	Longitude	House Price
MedInc	-0.079809	-0.015176	0.688075
HouseAge	0.011173	-0.108197	0.105623
AveRooms	0.106389	-0.027540	0.151948
AveBedrms	0.069721	0.013344	-0.046701
Population	-0.108785	0.099773	-0.024650
AveOccup	0.002366	0.002476	-0.023737
Latitude	1.000000	-0.924664	-0.144160
Longitude	-0.924664	1.000000	-0.045967
House Price	-0.144160	-0.045967	1.000000

```
import seaborn as sns
sns.heatmap(corr,annot=True,cmap='magma')
```

<Axes: >



```
#Feature Selection
import numpy as np
x=np.array(df)
Y=x[:,8]
x=x[:,0:7]

#Splitting the Dataset into Training and Testing
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(x,Y,random_state=42,test_size=0.25)
```

C. Linear Regression Model & E. Evaluation Metrics

```
#Fitting the Dataset into Liner Regression Model
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()
```



```

#Predicting the X_test
y_pred = model.predict(X_test)
print(y_pred)

[1.00061777 1.57882678 2.59058544 ... 1.76171598 2.76006792
3.60760481]

#Finding the r2 Score and Mean Squared Error
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared Score: {r2}")

Mean Squared Error: 0.6232181213603569
R-squared Score: 0.5290133054812776

```

D. Hyperparameter Tuning and Model Evaluation & E. Evaluation Metrics:

```

from sklearn.preprocessing import StandardScaler

# Standardization using StandardScaler
scaler_standard = StandardScaler()
standardized_data = scaler_standard.fit_transform(df)
standardized_df = pd.DataFrame(standardized_data, columns=df.columns)

model = LinearRegression(fit_intercept=False)
model.fit(X_train, y_train)

LinearRegression(fit_intercept=False)

#predicting the X_test
y_pred = model.predict(X_test)
print(y_pred)

[2.89636454 1.10469645 3.06203802 ... 1.6674325 2.6980249
1.87189299]


#Calculating and comparing the Root Mean Square Error (RMSE) and r2
Score with Hyperparameters
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared Score: {r2}")

Mean Squared Error: 0.666295686924439
R-squared Score: 0.5264305608038502

```

```
plt.scatter(y_test, y_pred,color='violet',s=10)
plt.plot([min(Y), max(Y)], [min(Y), max(Y)], color='black')
plt.xlabel('Actual y')
plt.ylabel('Predicted y')
plt.title('Actual vs Predicted y')
plt.show()
```



## INTERPRETATION

Accuracy without Hyperparameter :- 52% Accuracy with Hyperparameter :- 52%

The consistent R2 score between models with and without hyperparameter tuning suggests that either the default hyperparameters are already optimal for the dataset, the dataset itself is not highly complex, or the chosen hyperparameter range may not encompass the optimal values. It's crucial to consider expanding the hyperparameter search space or exploring alternative metrics to ensure thorough model optimization and performance evaluation.