

dgaypjm

April 4, 2024

Lab 06

```
[215]: # Name :- Arun M & Falak Ansari
# Reg No :- 23122110 & 23122106
# Date :- 04-04-2024
```

```
[216]: from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline,make_pipeline
from sklearn.feature_selection import SelectKBest,chi2
import pandas as pd
```

```
[217]: #Read the Smaller verison of the Dataset
df = pd.read_excel('copy.xlsx')
df
```

```
[217]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
5	6	0	3	
6	7	0	1	
7	8	0	3	
8	9	1	3	
9	10	1	2	
10	11	1	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	

4	Allen, Mr. William Henry	male	35.0	0
5	Moran, Mr. James	male	NaN	0
6	McCarthy, Mr. Timothy J	male	54.0	0
7	Palsson, Master. Gosta Leonard	male	2.0	3
8	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27.0	0
9	Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1
10	Sandstrom, Miss. Marguerite Rut	female	4.0	1

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
5	0	330877	8.4583	NaN	Q
6	0	17463	51.8625	E46	S
7	1	349909	21.0750	NaN	S
8	2	347742	11.1333	NaN	S
9	0	237736	30.0708	NaN	C
10	1	PP 9549	16.7000	G6	S

```
[218]: df.drop(columns=['PassengerId', 'Name', 'Ticket', 'Cabin'], inplace=True)
```

```
[219]: #Increased the test size of the dataset
X_train, X_test, y_train, y_test = train_test_split(df.
↳ drop(columns=['Survived']), df['Survived'], test_size=0.4, random_state=42)
```

```
[220]: X_train
```

```
[220]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
1	1	female	38.0	1	0	71.2833	C
8	3	female	27.0	0	2	11.1333	S
4	3	male	35.0	0	0	8.0500	S
7	3	male	2.0	3	1	21.0750	S
3	1	female	35.0	1	0	53.1000	S
6	1	male	54.0	0	0	51.8625	S

```
[221]: # imputation transformer
trf1 = ColumnTransformer([
    ('impute_age', SimpleImputer(), [2]),
    ('impute_embarked', SimpleImputer(strategy='most_frequent'), [6])
], remainder='passthrough')
# fit the preprocessor to the training data
encoded_data = trf1.fit_transform(X_train)

#Printing the encoded data
print(encoded_data)
```

```

[[38.0 'C' 1 'female' 1 0 71.2833]
 [27.0 'S' 3 'female' 0 2 11.1333]
 [35.0 'S' 3 'male' 0 0 8.05]
 [2.0 'S' 3 'male' 3 1 21.075]
 [35.0 'S' 1 'female' 1 0 53.1]
 [54.0 'S' 1 'male' 0 0 51.8625]]

```

```

[222]: # one hot encoding
trf2 = ColumnTransformer([
    ↳
    ↳↳('ohe_sex_embarked',OneHotEncoder(sparse_output=False,handle_unknown='ignore'),[1,6])
],remainder='passthrough')

#Applied the OneHotEncoder to the dataframe
encoded_data2=trf2.fit_transform(X_train)

#Printed the encoded data
print(encoded_data2)

```

```

[[ 1.      0.      1.      0.      1.      38.      1.      0.      71.2833]
 [ 1.      0.      0.      1.      3.      27.      0.      2.      11.1333]
 [ 0.      1.      0.      1.      3.      35.      0.      0.       8.05 ]
 [ 0.      1.      0.      1.      3.       2.      3.      1.      21.075 ]
 [ 1.      0.      0.      1.      1.      35.      1.      0.      53.1   ]
 [ 0.      1.      0.      1.      1.      54.      0.      0.      51.8625]]

```

```

[223]: # Scaling
trf3 = ColumnTransformer([
    ('scale',MinMaxScaler(),slice(0,10))
])

```

```

[224]: # Feature selection
trf4 = SelectKBest(score_func=chi2,k=8)

```

```

[225]: clf = LogisticRegression()

```

```

[226]: pipe = Pipeline([
    ('trf1',trf1),
    ('trf2',trf2),
    ('trf3',trf3),
    ('trf4',trf4),
    ('clf',clf)
])

```

```

[227]: # train
pipe.fit(X_train,y_train)

```

```
[227]: Pipeline(steps=[('trf1',
                        ColumnTransformer(remainder='passthrough',
                                           transformers=[('impute_age', SimpleImputer(),
                                                           [2]),
                                                           ('impute_embarked',
                                                           SimpleImputer(strategy='most_frequent'),
                                                           [6])])),
                      ('trf2',
                        ColumnTransformer(remainder='passthrough',
                                           transformers=[('ohe_sex_embarked',
                                                           OneHotEncoder(handle_unknown='ignore',
                                                           sparse_output=False),
                                                           [1, 6])])),
                      ('trf3',
                        ColumnTransformer(transformers=[('scale', MinMaxScaler(),
                                                           slice(0, 10, None))])),
                      ('trf4',
                        SelectKBest(k=8,
                                    score_func=<function chi2 at 0x000001A25B6A4C10>)),
                      ('clf', LogisticRegression())])
```

```
[228]: # Display Pipeline

from sklearn import set_config
set_config(display='diagram')
```

```
[229]: # Predict
y_pred = pipe.predict(X_test)
```

```
[230]: from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

[230]: 0.6

```
[231]: # cross validation using cross_val_score
from sklearn.model_selection import cross_val_score
cross_val_score(pipe, X_train, y_train, cv=5, scoring='accuracy').mean()
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[231], line 3
      1 # cross validation using cross_val_score
      2 from sklearn.model_selection import cross_val_score
----> 3 cross_val_score(pipe, X_train, y_train, cv=5, scoring='accuracy').mean()
```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection
↳py:562, in cross_val_score(estimator, X, y, groups, scoring, cv, n_jobs,
↳verbose, fit_params, pre_dispatch, error_score)
    559 # To ensure multimetric format is not supported
    560 scorer = check_scoring(estimator, scoring=scoring)
--> 562 cv_results = cross_validate(
    563     estimator=estimator,
    564     X=X,
    565     y=y,
    566     groups=groups,
    567     scoring={"score": scorer},
    568     cv=cv,
    569     n_jobs=n_jobs,
    570     verbose=verbose,
    571     fit_params=fit_params,
    572     pre_dispatch=pre_dispatch,
    573     error_score=error_score,
    574 )
    575 return cv_results["test_score"]

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\utils\_param_val
↳py:214, in validate_params.<locals>.decorator.<locals>.wrapper(*args, **kwargs)
    208 try:
    209     with config_context(
    210         skip_parameter_validation=(
    211             prefer_skip_nested_validation or global_skip_validation
    212         )
    213     ):
--> 214     return func(*args, **kwargs)
    215 except InvalidParameterError as e:
    216     # When the function is just a wrapper around an estimator, we allow
    217     # the function to delegate validation to the estimator, but we
↳replace
    218     # the name of the estimator by the name of the function in the error
    219     # message to avoid confusion.
    220     msg = re.sub(
    221         r"parameter of \w+ must be",
    222         f"parameter of {func.__qualname__} must be",
    223         str(e),
    224     )

```

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection
↳py:309, in cross_validate(estimator, X, y, groups, scoring, cv, n_jobs,
↳verbose, fit_params, pre_dispatch, return_train_score, return_estimator,
↳return_indices, error_score)
    306 # We clone the estimator to make sure that all the folds are

```

```

307 # independent, and that it is pickle-able.
308 parallel = Parallel(n_jobs=n_jobs, verbose=verbose,
↳pre_dispatch=pre_dispatch)
--> 309 results = parallel(
310     delayed(_fit_and_score)(
311         clone(estimator),
312         X,
313         y,
314         scorers,
315         train,
316         test,
317         verbose,
318         None,
319         fit_params,
320         return_train_score=return_train_score,
321         return_times=True,
322         return_estimator=return_estimator,
323         error_score=error_score,
324     )
325     for train, test in indices
326 )
328 _warn_or_raise_about_fit_failures(results, error_score)
330 # For callable scoring, the return type is only know after calling. If
↳the
331 # return type is a dictionary, the error scores can now be inserted with
332 # the correct key.

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\utils\parallel.
py:65, in Parallel.__call__(self, iterable)
60 config = get_config()
61 iterable_with_config = (
62     (_with_config(delayed_func, config), args, kwargs)
63     for delayed_func, args, kwargs in iterable
64 )
---> 65 return super().__call__(iterable_with_config)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\joblib\parallel.
py:1863, in Parallel.__call__(self, iterable)
1861 output = self._get_sequential_output(iterable)
1862 next(output)
-> 1863 return output if self.return_generator else list(output)
1865 # Let's create an ID that uniquely identifies the current call. If the
1866 # call is interrupted early and that the same instance is immediately
1867 # re-used, this id will be used to prevent workers that were
1868 # concurrently finalizing a task from the previous call to run the
1869 # callback.

```

```

1870 with self._lock:

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\joblib\parallel.
↳py:1789, in Parallel._get_sequential_output(self, iterable)
    1786 yield None
    1788 # Sequentially call the tasks and yield the results.
-> 1789 for func, args, kwargs in iterable:
    1790     self.n_dispatched_batches += 1
    1791     self.n_dispatched_tasks += 1

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\utils\parallel.
↳py:61, in <genexpr>(.0)
    56 # Capture the thread-local scikit-learn configuration at the time
    57 # Parallel.__call__ is issued since the tasks can be dispatched
    58 # in a different thread depending on the backend and on the value of
    59 # pre_dispatch and n_jobs.
    60 config = get_config()
---> 61 iterable_with_config = (
    62     (_with_config(delayed_func, config), args, kwargs)
    63     for delayed_func, args, kwargs in iterable
    64 )
    65 return super().__call__(iterable_with_config)

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection.
↳py:309, in <genexpr>(.0)
    306 # We clone the estimator to make sure that all the folds are
    307 # independent, and that it is pickle-able.
    308 parallel = Parallel(n_jobs=n_jobs, verbose=verbose,
↳pre_dispatch=pre_dispatch)
--> 309 results = parallel(
    310     delayed(_fit_and_score)(
    311         clone(estimator),
    312         X,
    313         y,
    314         scorers,
    315         train,
    316         test,
    317         verbose,
    318         None,
    319         fit_params,
    320         return_train_score=return_train_score,
    321         return_times=True,
    322         return_estimator=return_estimator,
    323         error_score=error_score,
    324     )

```

```

325     for train, test in indices
326 )
328 _warn_or_raise_about_fit_failures(results, error_score)
330 # For callable scoring, the return type is only know after calling. If
↳the
331 # return type is a dictionary, the error scores can now be inserted with
332 # the correct key.

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection\
↳py:377, in _BaseKFold.split(self, X, y, groups)
369 if self.n_splits > n_samples:
370     raise ValueError(
371         (
372             "Cannot have number of splits n_splits={0} greater"
373             " than the number of samples: n_samples={1}."
374         ).format(self.n_splits, n_samples)
375     )
--> 377 for train, test in super().split(X, y, groups):
378     yield train, test

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection\
↳py:108, in BaseCrossValidator.split(self, X, y, groups)
106 X, y, groups = indexable(X, y, groups)
107 indices = np.arange(_num_samples(X))
--> 108 for test_index in self._iter_test_masks(X, y, groups):
109     train_index = indices[np.logical_not(test_index)]
110     test_index = indices[test_index]

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection\
↳py:770, in StratifiedKFold._iter_test_masks(self, X, y, groups)
769 def _iter_test_masks(self, X, y=None, groups=None):
--> 770     test_folds = self._make_test_folds(X, y)
771     for i in range(self.n_splits):
772         yield test_folds == i

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection\
↳py:732, in StratifiedKFold._make_test_folds(self, X, y)
730 min_groups = np.min(y_counts)
731 if np.all(self.n_splits > y_counts):
--> 732     raise ValueError(
733         "n_splits=%d cannot be greater than the"
734         " number of members in each class." % (self.n_splits)
735     )
736 if self.n_splits > min_groups:
737     warnings.warn(

```



```

738         "The least populated class in y has only %d"
739         " members, which is less than n_splits=%d."
740         % (min_groups, self.n_splits),
741         UserWarning,
742     )

```

ValueError: n_splits=5 cannot be greater than the number of members in each class.

```

[ ]: # gridsearchcv
param_grid = {
    "clf__penalty": ['l1', 'l2'],
    "clf__C": [0.001, 0.01, 0.1, 1, 10, 100]
}

```

```

[ ]: from sklearn.model_selection import GridSearchCV
grid = GridSearchCV(pipe, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train, y_train)

```

C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection_split.py:737: UserWarning: The least populated class in y has only 3 members, which is less than n_splits=5.

warnings.warn(

C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection_validation.py:425: FitFailedWarning: 36 fits failed out of a total of 60.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

24 fits failed with the following error:

Traceback (most recent call last):

File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection_validation.py", line 729, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\base.py", line 1152, in wrapper
return fit_method(estimator, *args, **kwargs)

File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-

```

packages\sklearn\pipeline.py", line 427, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\base.py", line 1152, in wrapper
    return fit_method(estimator, *args, **kwargs)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\linear_model\_logistic.py", line 1169, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\linear_model\_logistic.py", line 56, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

```

12 fits failed with the following error:

Traceback (most recent call last):

```

File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\base.py", line 1152, in wrapper
    return fit_method(estimator, *args, **kwargs)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\pipeline.py", line 423, in fit
    Xt = self._fit(X, y, **fit_params_steps)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\pipeline.py", line 377, in _fit
    X, fitted_transformer = fit_transform_one_cached(
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\joblib\memory.py", line 353, in __call__
    return self.func(*args, **kwargs)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\pipeline.py", line 957, in _fit_transform_one
    res = transformer.fit_transform(X, y, **fit_params)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\utils\_set_output.py", line 157, in wrapped
    data_to_wrap = f(self, X, *args, **kwargs)
File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

```

10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\base.py", line 1152, in wrapper
    return fit_method(estimator, *args, **kwargs)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\compose\_column_transformer.py", line 754, in fit_transform
    result = self._fit_transform(X, y, _fit_transform_one)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\compose\_column_transformer.py", line 681, in _fit_transform
    return Parallel(n_jobs=self.n_jobs)(
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\utils\parallel.py", line 65, in __call__
    return super().__call__(iterable_with_config)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\joblib\parallel.py", line 1863, in __call__
    return output if self.return_generator else list(output)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\joblib\parallel.py", line 1792, in _get_sequential_output
    res = func(*args, **kwargs)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\utils\parallel.py", line 127, in __call__
    return self.function(*args, **kwargs)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\pipeline.py", line 957, in _fit_transform_one
    res = transformer.fit_transform(X, y, **fit_params)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\utils\_set_output.py", line 157, in wrapped
    data_to_wrap = f(self, X, *args, **kwargs)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\base.py", line 919, in fit_transform
    return self.fit(X, y, **fit_params).transform(X)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\preprocessing\_data.py", line 435, in fit
    return self.partial_fit(X, y)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\base.py", line 1152, in wrapper
    return fit_method(estimator, *args, **kwargs)
    File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

```

10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\preprocessing\_data.py", line 473, in partial_fit
    X = self._validate_data(
      File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\base.py", line 605, in _validate_data
        out = check_array(X, input_name="X", **check_params)
      File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\utils\validation.py", line 915, in check_array
        array = _asarray_with_order(array, order=order, dtype=dtype, xp=xp)
      File "C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\utils\_array_api.py", line 380, in _asarray_with_order
        array = numpy.asarray(array, order=order, dtype=dtype)
ValueError: could not convert string to float: 'female'

```

```

warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\arunp\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.10_qbz5n
2kfra8p0\LocalCache\local-packages\Python310\site-
packages\sklearn\model_selection\_search.py:979: UserWarning: One or more of the
test scores are non-finite: [nan nan nan nan nan nan nan nan nan nan]
    warnings.warn(

```

ValueError Traceback (most recent call last)

Cell In[56], line 3

```

1 from sklearn.model_selection import GridSearchCV
2 grid = GridSearchCV(pipe, param_grid, cv=5, scoring='accuracy')
----> 3 grid.fit(X_train, y_train)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↪10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\base.
↪py:1152, in _fit_context.<locals>.decorator.<locals>.wrapper(estimator, *args,
↪**kwargs)

```

```

    1145 estimator._validate_params()
    1147 with config_context(
    1148     skip_parameter_validation=(
    1149         prefer_skip_nested_validation or global_skip_validation
    1150     )
    1151 ):
-> 1152     return fit_method(estimator, *args, **kwargs)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↪10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\model_selection\
↪py:936, in BaseSearchCV.fit(self, X, y, groups, **fit_params)
    934 refit_start_time = time.time()
    935 if y is not None:

```

```

--> 936     self.best_estimator_.fit(X, y, **fit_params)
      937 else:
      938     self.best_estimator_.fit(X, **fit_params)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↪10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\base.
↪py:1152, in _fit_context.<locals>.decorator.<locals>.wrapper(estimator, *args,
↪**kwargs)
      1145     estimator._validate_params()
      1147 with config_context(
      1148     skip_parameter_validation=(
      1149         prefer_skip_nested_validation or global_skip_validation
      1150     )
      1151 ):
--> 1152     return fit_method(estimator, *args, **kwargs)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↪10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\pipeline.
↪py:427, in Pipeline.fit(self, X, y, **fit_params)
      425     if self._final_estimator != "passthrough":
      426         fit_params_last_step = fit_params_steps[self.steps[-1][0]]
--> 427         self._final_estimator.fit(Xt, y, **fit_params_last_step)
      429 return self

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↪10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\base.
↪py:1152, in _fit_context.<locals>.decorator.<locals>.wrapper(estimator, *args,
↪**kwargs)
      1145     estimator._validate_params()
      1147 with config_context(
      1148     skip_parameter_validation=(
      1149         prefer_skip_nested_validation or global_skip_validation
      1150     )
      1151 ):
--> 1152     return fit_method(estimator, *args, **kwargs)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.

```

↪10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\linear_model\_l
↪py:1169, in LogisticRegression.fit(self, X, y, sample_weight)
      1139 @_fit_context(prefer_skip_nested_validation=True)
      1140 def fit(self, X, y, sample_weight=None):
      1141     """
      1142     Fit the model according to the given training data.
      1143
      1144     (...)
      1167     The SAGA solver supports both float64 and float32 bit arrays.
      1168     """
--> 1169     solver = _check_solver(self.solver, self.penalty, self.dual)
      1171     if self.penalty != "elasticnet" and self.l1_ratio is not None:

```

```

1172         warnings.warn(
1173             "l1_ratio parameter is only used when penalty is "
1174             "'elasticnet'. Got "
1175             "(penalty={})".format(self.penalty)
1176         )

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.
↳10_qbz5n2kfra8p0\LocalCache\local-packages\Python310\site-packages\sklearn\linear_model\_l
↳py:56, in _check_solver(solver, penalty, dual)
    53 def _check_solver(solver, penalty, dual):
    54     # TODO(1.4): Remove "none" option
    55     if solver not in ["liblinear", "saga"] and penalty not in ("l2",
↳"none", None):
--> 56         raise ValueError(
    57             "Solver %s supports only 'l2' or 'none' penalties, got %s"
↳penalty."
    58             % (solver, penalty)
    59         )
    60     if solver != "liblinear" and dual:
    61         raise ValueError(
    62             "Solver %s supports only dual=False, got dual=%s" % (solver
↳dual)
    63         )

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty

```

```
[ ]: grid.best_score_
```

```
[ ]: nan
```

```
[ ]: grid.best_params_
```

```
[ ]: {'clf__C': 0.001, 'clf__penalty': 'l1'}
```

```
[ ]: # export
import pickle
pickle.dump(pipe, open('lr_pile.pkl', 'wb'))
```